

SNEIKEN and SNEIKHA

Authenticated Encryption and Cryptographic Hashing (Preliminary version of Wednesday 27th February, 2019)

Markku-Juhani O. Saarinen

PQShield Ltd.
Prima House, 267 Banbury Road
Oxford OX2 7HQ, United Kingdom
mjos@pqshield.com
Tel. +44 (0)7548 620723

Abstract. We describe the lightweight SNEIK permutation and two derived sponge modes: the SNEIKEN Authenticated Encryption with Associated Data (AEAD) construction and the SNEIKHA Cryptographic Hash. The permutation is a simple ARX design with very efficient feedback mixing, and is optimized for low-end microcontrollers. The overall design emphasizes simplicity, small implementation footprint, and ease of integration with lightweight cryptographic protocols and post-quantum schemes. The submission package includes implementations for Atmel AVR and ARM Cortex M3/M4 targets, where SNEIK performs better than comparable permutations and AES. However, we see small RAM and ROM implementation footprint as the main advantage of SNEIK and reserve the right to double the number of rounds if there is progress in its cryptanalysis.

Keywords: Lightweight Cryptography · Sponge Modes · SNEIKEN · SNEIKHA

Contents

1	Introduction	2
2	The SNEIK Permutation	3
3	BLNK2 Primitive Sponge Operations	4
4	The SNEIKEN Authenticated Encryption Algorithm	7
5	The SNEIKHA Cryptographic Hash	8
6	Design Rationale	9



Copyright © 2019 PQShield Ltd., Oxford UK.

1 Introduction

This document describes the SNEIK family of primitives for lightweight cryptography. The primary members of the family are the **SNEIKEN128** AEAD (Authenticated Encryption with Associated Data) algorithm and the **SNEIKHA256** cryptographic hash. SNEIKEN256 and SNEIKHA384 can be paired for higher-security applications.

Name	Type	Security	Specification
SNEIKEN128	AEAD	2^{128}	Section 4.
SNEIKEN256	AEAD	2^{256}	Section 4.
<i>SNEIQEN128</i>	AEAD		Section 4.
SNEIKHA256	HASH	2^{128}	Section 5.
SNEIKHA384	HASH	2^{192}	Section 5.
<i>SNEIGEN128</i>	XOF		Section 5

The classical security for (SNEIKEN) AEADs indicates the effort required to breach the confidentiality of a given plaintext with a classical computer, and is equivalent to key size. The effort required to breach integrity of ciphertext (i.e. to create a forgery) is claimed to be equivalent to size of the ciphertext expansion (authentication tag). Any valid attack must ensure that a nonce does not repeat under the same secret key.

For (SNEIKHA) hash functions we primarily indicate the effort required to produce collisions on a classical computer. (Second) pre-image attacks may require more effort, especially for fixed-format or short messages, as used in some hash-based signatures.

We set no explicit limits on the input sizes (hashed message, plaintext, associated data, and the amount of data that can be processed under one key), but we assume it to be under 2^{64} bits for security analysis.

The SNEIQEN128 AEAD and SNEIGEN128 XOF are included as “informational”. Even though they have clear use cases in lightweight cryptography, they may not meet the most stringent security criteria for all applications. They are intended as “building blocks” instead; their security must be evaluated in the context where they are used.

Shared features between AEAD and Hash. The SNEIKEN and SNEIKHA proposals share the underlying SNEIK permutation f_{512}^p (Section 2), and the BLNK2 padding mechanism (Section 3). Implementations of the two algorithms may have up to 90% common code, as can be seen from the reference implementations provided.

We note that SNEIK is intended as a fully-featured suite that fulfills all symmetric cryptographic needs of a lightweight application. The BLNK2 modes are based on Author’s BLINKER framework for Sponge-based protocols [Saa14a], which has inspired derivative works such as Mike Hamburg’s lightweight STROBE protocol [Ham17].

Notation and conventions. SNEIK is an ARX [KN10] type construction built from three very simple operations on 32-bit words:

- A: $x \boxplus y$ Addition modulo word size: $x + y \bmod 2^{32}$.
- R: $x \oplus y$ Bitwise exclusive-or operation between x and y .
- X: $x \lll r$ Cyclic left rotation by r bits in 32-bit word.

We also use Boolean operators \wedge and \vee to denote bitwise “and” and “or” operations and vertical $\|$ to denote concatenation of arrays and strings.

C-style notation is used for bit and byte arrays; vectors are zero-indexed with index in square brackets. We use ranges to indicate sub-arrays; $v[i \dots j]$ refers to concatenation of all entries from $v[i]$ to $v[j]$, inclusive.

All numerical values are stored and exchanged in little-endian fashion, with the least significant bit, byte, or vector array entry having index 0. Hexadecimal numbers (bytes or

```

// cyclic rotate left for 32-bit words
#define ROL32(x, y) (((x) << (y)) | ((x) >> (32 - (y))))

void sneik_f512(void *s, uint8_t dom, uint8_t rounds)
{
    const uint8_t rc[16] = {
        0xEF, 0xE0, 0xD9, 0xD6, 0xBA, 0xB5, 0x8C, 0x83,
        0x10, 0x1F, 0x26, 0x29, 0x45, 0x4A, 0x73, 0x7C // (only 8 used now)
    };

    int i, j;
    uint32_t t, *v = (uint32_t *) s; // loop counters
    // assume little endian!

    for (i = 0; i < rounds; i++) { // loop over rounds
        v[0] ^= (uint32_t) rc[i]; // xor round constant
        v[1] ^= (uint32_t) dom; // xor domain constant
        for (j = 0; j < 16; j++) {
            t = v[j]; // middle value
            t += v[(j - 1) & 0xF]; // feedback previous
            t = t ^ ROL32(t, 24) ^ ROL32(t, 25); // p(x) = x^25 + x^24 + x
            t ^= v[(j - 2) & 0xF]; // outer feedback
            t += v[(j + 2) & 0xF];
            t = t ^ ROL32(t, 9) ^ ROL32(t, 17); // q(x) = x^17 + x^9 + x
            t ^= v[(j + 1) & 0xF]; // reverse feedback
            v[j] = t; // store the result
        }
    }
}

```

Listing 1: The SNEIK permutation $\text{f512}_\delta^\rho(S)$ in C. We set $\text{dom} = \delta$ and $\text{rounds} = \rho$.

words) are prefixed with “0x”. Bit and byte arrays are read from left to right, with index starting with 0. The 32-bit integer 0x12345678 (decimal 305419896) is therefore stored and transmitted as four bytes 0x78 || 0x56 || 0x34 || 0x12.

Any integer $n \in (0, 2^m]$ has unique encoding as bit array $B[m]$ with $n = \sum_i^{m-1} 2^i B[i]$. Therefore bit i has numerical value 2^i and the first bit (bit 0) of a byte is 0x01 and the last bit (bit 7) is $2^7 = 0x80$. One can always fetch bit i from a byte array $v[]$ in C with an expression such as $(v[i >> 3] >> (i \& 7)) \& 1$.

2 The SNEIK Permutation

With π_δ^ρ we denote a family of ρ -round permutations on b -bit state S , controlled by a domain identifier δ :

$$S' = \pi_\delta^\rho(S). \quad (1)$$

Listing 1 contains a compact C source code implementation of the SNEIK permutation instantiation $\pi = \text{f512}$ (with $b = 512$) used in our SNEIKEN and SNEIKHA proposals.

Non-linear feedback shift register. Let $n \geq 5$ be the size of the initial state $s[0 \dots n-1]$ of 32-bit words (with the f512 instantiation we have $n = 16$). Recurrence of Equation 2 defines a nonlinear feedback expander sequence $s[i]$ for $i \geq n$. The seven arithmetic steps

Table 1: SNEIK permutation performance on 32-bit ARM Cortex-M4 (NXP/Freescale MK20DX256) and 8-bit AVR (Atmel ATMEGA2560) architectures. The “RAM” size is the input/output state + stack usage. Cycles per round was measured with $\rho = 8$.

MCU	Unroll	RAM	ROM	Cycles/Round
AVR	16-step	64 + 14	1974	1078.1
AVR	4-step	64 + 19	618	1126.0
Cortex M4	16-step	64 + 16	560	188.0
Cortex M4	4-step	64 + 28	232	211.8

t_j are numbered just for referencing.

$$\begin{aligned}
t_1 &= s[i - n] \oplus d[i] \\
t_2 &= t_1 \boxplus s[i - 1] \\
t_3 &= t_2 \oplus (t_2 \lll 24) \oplus (t_2 \lll 25) \\
t_4 &= t_3 \oplus s[i - 2] \\
t_5 &= t_4 \boxplus s[i - n + 2] \\
t_6 &= t_5 \oplus (t_5 \lll 9) \oplus (t_5 \lll 17) \\
t_7 &= t_6 \oplus s[i - n + 1] \\
s[i] &= t_7
\end{aligned} \tag{2}$$

The domain separation constant $d[i]$ is nonzero only when $i \bmod n \in \{0, 1\}$. We interpret round constants to be just another kind of “domain separator”, separating rounds from each other. We set $d[nj] = rc[j]$ from vector in Equation 3 and $d[nj + 1] = \delta$. The domain identifier value of δ is set by higher level primitive (see Table 2).

$$\begin{aligned}
rc[0..15] &= 0xEF, 0xE0, 0xD9, 0xD6, 0xBA, 0xB5, 0x8C, 0x83, \\
&0x10, 0x1F, 0x26, 0x29, 0x45, 0x4A, 0x73, 0x7C
\end{aligned} \tag{3}$$

Implementation: “Sliding window”. Since there are no references beyond $s[i - n]$ back in the sequence, the recurrence of Equation 2 may be implemented with a static n -word table – as was done in Listing 1. We may use mod n “addressing” and write $s[i - n \pm j]$ as $s[i \pm j]$ while i repeatedly scans the values $i = 0, 1, \dots, n - 1$ for each round.

We see that the operation uses a “window” of five inputs to evaluate each new value:

$$s[i] = f_{\text{win}}(s[i - 2], s[i - 1], s[i], s[i + 1], s[i + 2]) \tag{4}$$

Four 32-bit state words can be used to store the f inputs as the window moves; the value $s[i - 2]$ is used at step t_4 before a replacement value $s[i + 2]$ is loaded for step t_5 .

The standard implementation method is therefore to unroll computation of at least four iterations of Equation 2. Table 1 gives some implementation metrics for the permutation on popular microcontrollers using this method.

3 BLNK2 Primitive Sponge Operations

Our proposals are built from lower-level “BLINKER-style” [Saa14a] primitives. In addition to authenticated encryption and hashing, these primitives can be used to build more complex protocols where two (or more) parties have synchronized, continuously authenticated states. We write the block size as $b = 512$ and the rate is $r = 128/256/384$.

For these modes a tuple (S, i) defines the entire state: $S \in \{0, 1\}^b$ is the permutation input/output and $i \in [0, b)$ is a “next bit” read/write index to it at $S[i]$. The primitives may set additional flags on domain parameter δ before passing them to the cryptographic permutation π_δ^ρ . This 8-bit domain identifier is constructed from fields given in Table 2.

$S.\text{clr}()$	Clear the state: $S \leftarrow 0^b, i \leftarrow 0$.
$S.\text{fin}(\delta)$	Mark the end of given domain (Algorithm 2).
$S.\text{put}(D, \delta)$	Absorb input data D (Algorithm 3).
$D \leftarrow S.\text{get}(n, \delta)$	Squeeze out n bits into D (Algorithm 4).
$C \leftarrow S.\text{enc}(P, \delta)$	Encrypt plaintext P into ciphertext C (Algorithm 5).
$P \leftarrow S.\text{dec}(C, \delta,)$	Decrypt ciphertext C into plaintext P (Algorithm 6).

Additionally, we have a utility function $S.\text{inc}(\delta)$ (Algorithm 1) which updates the index i by one and invokes the permutation π_δ^ρ if it reaches the limit set by rate r or block b , depending on the full bit in the domain indicator δ .

Algorithm 1 Increment index: $S.\text{inc}(\delta)$.

Input: Input state (S, i) , domain δ

- | | |
|---|--|
| 1: $i \leftarrow i + 1$ | <i>Increment index.</i> |
| 2: if $(\delta \wedge \text{full} = 0 \text{ and } i = r)$ or
$(\delta \wedge \text{full} = \text{full} \text{ and } i = b)$ then | |
| 3: $S \leftarrow \pi_\delta^\rho(S)$ | <i>Apply permutation if rate or block is full.</i> |
| 4: $i \leftarrow 0$ | <i>Reset index.</i> |
| 5: end if | |

Output: Updated state (S, i) .

Algorithm 2 End a data element (padding): $S.\text{fin}(\delta)$.

Input: Input state (S, i) , domain δ

- | | |
|---|---|
| 1: $S[i] \leftarrow S[i] \oplus 1$ | <i>Add padding bit, typically byte 0x01.</i> |
| 2: if $\delta \wedge \text{full} = 0$ then | |
| 3: $S[r - 1] \leftarrow S[r - 1] \oplus 1$ | <i>Normal capacity; last rate byte gets 0x80.</i> |
| 4: end if | |
| 5: $S \leftarrow \pi_{(\delta \vee \text{last})}^\rho(S)$ | <i>Permutation with domain end marker last.</i> |
| 6: $i \leftarrow 0$ | <i>Reset index.</i> |

Output: Updated state (S, i) .

Algorithm 3 Absorb data: $S.\text{put}(D, \delta)$.

Input: Input state (S, i) , data $D \in \{0, 1\}^*$, domain δ .

- | | |
|---|---|
| 1: for $j = 0, 1, \dots, \text{length}(D) - 1$ do | |
| 2: $S[i] \leftarrow S[i] \oplus D[j]$ | <i>Add (xor) input data to the state.</i> |
| 3: $S.\text{inc}(\delta)$ | <i>Increment index i.</i> |
| 4: end for | |

Output: Updated state (S, i) .

Algorithm 4 Squeeze data: $D = S.get(n, \delta)$.

Input: Input state (S, i) , length of output n , domain δ .

```

1:  $D = \{\}$  Empty string.
2: for  $j = 0, 1, \dots, n - 1$  do
3:    $D[j] \leftarrow S[i]$  Get a bit from the state.
4:    $S.inc(\delta)$  Increment index  $i$ .
5: end for
```

Output: Output data D , updated state (S, i) .

Algorithm 5 Encrypt data: $C = S.enc(P, \delta)$.

Input: Input state (S, i) , plaintext P , domain δ .

```

1:  $C = \{\}$  Empty ciphertext.
2: for  $j = 0, 1, \dots, \text{length}(n)(P) - 1$  do
3:    $C[j] \leftarrow S[i] \oplus P[j]$  Xor plaintext with the state.
4:    $S[i] \leftarrow C[j]$  Ciphertext goes into the state.
5:    $S.inc(\delta)$  Increment index  $i$ .
6: end for
```

Output: Ciphertext C , updated state (S, i) .

Algorithm 6 Decrypt data: $P = S.dec(C, \delta)$.

Input: Input state (S, i) , ciphertext C , domain δ .

```

1:  $P = \{\}$  Empty plaintext.
2: for  $j = 0, 1, \dots, \text{length}(n)(P) - 1$  do
3:    $P[j] \leftarrow S[i] \oplus C[j]$  Xor ciphertext with the state.
4:    $S[i] \leftarrow C[j]$  Ciphertext goes into the state.
5:    $S.inc(\delta)$  Increment index  $i$ .
6: end for
```

Output: Plaintext P , updated state (S, i) .

Table 2: Domain indicator δ bits and fields.

Name	Value	Class	Purpose
last	0x01	Flag	Final (padded) block marker.
full	0x02	Flag	Full state indicator.
ad	0x10	Input	Authenticated Data / Hash input.
adf	0x12	Input	Full-state AAD ($adf = ad \vee full$).
key	0x20	Input	Secret key material.
keyf	0x22	Input	Initialization block ($keyf = key \vee full$).
hash	0x40	Output	Hash, MAC, or XOF.
ptct	0x70	In/out	Plaintext/ciphertext duplex block.

4 The SNEIKEN Authenticated Encryption Algorithm

The SNEIKEN family of authenticated encryption with associated data (AEAD) algorithms is characterized by the following six variables:

Var	Description	Length
K	Secret key	Fixed k
N	Nonce or IV	Fixed n
A	Associated data	Any a
P	Plaintext	Any p
T	Authentication tag	Fixed t
C	Ciphertext	$c = p + t$

The algorithms aim to provide integrity and confidentiality protection for P and C but only integrity protection for A . Capacity $c = b - r$ is equivalent to the key size k in encryption and decryption. Associated data is processed at full-state rate ($r = b$). Generally speaking, the confidentiality is at k -bit security level and integrity is at t -bit level (this may not hold for SNEIQEN128 in all attack models).

SNEIKEN128 is the primary member of the family:

Name	Rate	Rounds	Key	Nonce	Tag
SNEIKEN128	$r = 384$	$\rho = 6$	$k = 128$	$n = 128$	$t = 128$
SNEIKEN256	$r = 256$	$\rho = 8$	$k = 256$	$n = 128$	$t = 128$
SNEIQEN128	$r = 384$	$\rho = 4$	$k = 128$	$n = 96$	$t = 128$

Encryption and decryption. We define a 6-byte “variant identifier block” as follows:

$$\text{ID}[0..5] = 0\text{x}61, 0\text{x}65, r/8, k/8, n/8, t/8 \quad (5)$$

The first two bytes are ASCII ‘a’ and ‘e’, followed by byte lengths for rate, key, nonce, and tag. We denote the encryption process by $C \leftarrow \text{SNEIKEN}(K, N, A, P)$. Algorithm 7 contains the full procedure for SNEIKEN using the BLNK2 primitives defined in Section 3.

Algorithm 7 Authenticated encryption $C \leftarrow \text{SNEIKEN}(K, N, A, P)$.

Input: Secret key K , (public) nonce N , associated data A , and plaintext P .

- | | |
|---|--|
| 1: S.clr() | <i>Initialize the state: $S = 0^b, i = 0$</i> |
| 2: S.put(ID K N , keyf) | <i>Identifier, secret key, and nonce.</i> |
| 3: S.fin(keyf) | <i>Pad and permute the key block.</i> |
| 4: S.put(A , adf) | <i>Associated authenticated data.</i> |
| 5: S.fin(adf) | <i>Pad and permute, even if $a = 0$.</i> |
| 6: $C' \leftarrow \text{S.enc}(P, \text{ptct})$ | <i>Actual ciphertext.</i> |
| 7: S.fin(ptct) | <i>Pad and permute, even if $p = 0$.</i> |
| 8: $T \leftarrow \text{S.get}(t, \text{hash})$ | <i>Authentication tag, t bits.</i> |
| 9: $C \leftarrow C' T$ | <i>Authenticated ciphertext.</i> |

Output: Ciphertext C .

Algorithm 8 specifies the corresponding decryption and authentication function

$$\{P, \text{FAIL}\} \leftarrow \text{SNEIKEN}^{-1}(K, N, A, C). \quad (6)$$

Decryption must output only FAIL upon integrity check failure (no partial plaintext!)

Algorithm 8 Authenticated decryption $\{P, \text{FAIL}\} \leftarrow \text{SNEIKEN}^{-1}(K, N, A, C)$.

Input: Secret key K , (public) nonce N , associated data A , and ciphertext C .

1: $S.\text{clr}()$	<i>Initialize the state: $S = 0^b, i = 0$</i>
2: $S.\text{put}(\text{ID} \parallel K \parallel N, \text{keyf})$	<i>Identifier, secret key, and nonce.</i>
3: $S.\text{fin}(\text{keyf})$	<i>Pad and permute the key block.</i>
4: $S.\text{put}(A, \text{adf})$	<i>Associated authenticated data.</i>
5: $S.\text{fin}(\text{adf})$	<i>Pad and permute, even if $a = 0$.</i>
6: $P \leftarrow S.\text{dec}(C[0 \dots c - t - 1], \text{ptct})$	<i>Decrypt plaintext from first $c - t$ bits of C.</i>
7: $S.\text{fin}(\text{ptct})$	<i>Pad and permute, even if $p = 0$.</i>
8: $T = S.\text{get}(t, \text{hash})$	<i>Authentication tag, t bits.</i>
9: if $T = C[c - t \dots c - 1]$ then	
10: return P	<i>Last t bits of C matches with tag T.</i>
11: else	
12: return FAIL	<i>Authentication failure.</i>
13: end if	

Output: Plaintext P or FAIL.

Code Size. Compiling size-optimized `encrypt.c` that implements the NIST AEAD API (for Encryption and Decryption) resulted in 1100 bytes of executable code and data on AVR and 626 bytes on Cortex-M4. This is the only component required for implementation in addition to the permutation (Table 1). Full assembler implementation or co-implementation with SNEIKHA may yield smaller code size.

MAC-and-continue in lightweight setting. Lightweight protocols can avoid per-message rekeying by padding the MAC with $S.\text{fin}(\text{hash})$, and then directly continuing to process the next message (from step 4 in Algorithm 7). The decryption side must of course do the same. This is not only a significant speedup but also saves memory and provides “forward security” since there is no longer any need to retain the original secret key or nonce.

SNEIQEN Use Cases. The 4-round SNEIQEN may not be suitable as universally as the main SNEIKEN algorithms. It is intended for applications where an attacker has only a limited ability to perform chosen plaintext- or ciphertext queries – which is often the case with low-bandwidth and lightweight devices. The suitability of SNEIQEN must be evaluated individually for each application.

5 The SNEIKHA Cryptographic Hash

The SNEIKHA family of hash functions produce a h -bit hash H from input data A of arbitrary bit length a . The security against collision search for SNEIKHA algorithms is expected to be $2^{\frac{b-r}{2}}$ – which is equivalent to $2^{h/2}$ for these fixed-length hashes. Complexity of (second) pre-image search may be higher for format-restricted inputs.

SNEIKHA256 is the primary member of the family:

Name	Hash	Rate	Rounds	Security
SNEIKHA256	$h = 256$	$r = 256$	$\rho = 8$	2^{128}
SNEIKHA384	$h = 384$	$r = 128$	$\rho = 8$	2^{192}
<i>SNEIQEN128</i>	$h = \text{any}$	$r = 384$	$\rho = 4$	

Algorithm 9 specifies SNEIKHA using the BLNK2 primitives of Section 3. We note that if the squeezing step $S.\text{get}()$ is implemented literally (as in Algorithm 4), there may

be a final permutation call which is unnecessary if SNEIKHA is not used as a part of some intermediate-hash scheme. This is because, internally, the SNEIKHA algorithms are really extensible-output functions (XOFs). We may define explicit XOF padding modes in the future if a need arises to distinguish XOF use cases from fixed-length hashes.

Algorithm 9 Cryptographic hash $H \leftarrow \text{SNEIKHA}(A)$.

Input: Data to be hashed A .

1: $S.\text{clr}()$	<i>Initialize the state: $S = 0^b, i = 0$</i>
2: $S.\text{put}(A, \text{adf})$	<i>Absorb input data.</i>
3: $A \leftarrow S.\text{get}(h, \text{hash})$	<i>Squeeze hash, h bits.</i>

Output: Hash H of A .

Code Size. The size-optimized `hash.c` file implementing the NIST hash API compiles into 288 bytes on AVR and 180 bytes on Cortex-M4. This is the only component required for implementation in addition to the permutation (Table 1). Full assembler implementation or co-implementation with SNEIKEN may yield smaller code size. Incremental and keyed hashing constructions are straightforward.

SNEIGEN Use Cases. We also include SNEIGEN, which is really not a hash function but a seed expander with limited cryptographic strength. It is intended for cryptographic applications that need “random-like stuffing”. One such example is the padding in PKCS #1 [MKJR16]. Another example is the expansion of a short seed into public value \mathbf{A} in many lattice-based public key algorithms, including Round5 [BBF⁺19]. The authors of [BFM⁺18] argue that “good statistical properties” are sufficient for the public matrix \mathbf{A} in a lightweight implementation of the Frodo PQC encryption algorithm.

If the SNEIK permutation is used to build a general-purpose random number generator, this is also called “SNEIGEN”. New randomness can be added at any point with `S.put()`. If cryptographic security is required from the generator, we suggest increasing the number of rounds to $\rho = 8$ or even $\rho = 16$ and limiting rate to $r \leq b/2$.

6 Design Rationale

Design goals. Our main design goal was to create fast permutation-based primitives suitable for prominent 8, 16, and 32-bit embedded microcontrollers – primarily ARM Cortex-M and Atmel AVR families. The 32-bit Cortex-M target directly led to the use of a 32-bit primary datapath, while AVR somewhat limited the use of rotations (which are essentially “free” in Cortex M3/4).

We note that the size of the permutation n is actually entirely flexible – smaller and larger permutations can be easily constructed. This was one of the original design goals, although it is not used in the current proposals. However, it was clear that the entire permutation state would not fit into the register file of either of the main target platforms, so processing would have to be “localized” to some degree. This led to the “window” design of Equation 4. This is quite different from proposals such as Gimli [BKL⁺17], whose designers chose to have more localized mixing.

It was clear that the design should not have any table lookups or conditional branches, in order to make it naturally resistant to timing attacks and some other simple side-channel attacks. We toyed for a while with designs inspired by Ascon [DEMS16] (and therefore by Keccak and Xoodoo), but the fact that addition is “free” on the main target platforms finally made the decision to use ARX an easy one. NSA’s SPECK [BSS⁺13] was a strong inspiration in this sense. The overall structure is clearly influenced by a large number of

previous proposals, starting with the “Block TEA” algorithm by Wheeler and Needham (which the author cryptanalyzed more than two decades ago [WN98].)

XXX Shift-register – hardware implementation

XXX Extremely low state size – compare with Ascon, Xoofff

Strong feedback for fast avalanche. Since multiple-issue or superscalar processing is generally not available on lightweight targets, instruction and data path parallelism was not a great concern. Indeed, we decided to take an opposite route and maximize the critical path instead of minimizing it. As a result, we use immediate feedback from one processed word to the next, which helps to diffuse the state extremely rapidly. The design achieves complete avalanche (each input bit affecting each output bit) in only two rounds.

Round structure. The security of SNEIK relies largely on very effective feedback diffusion when the permutation is computed in either direction.

It is easy to see that each step in Equation 2 is invertible. The weight-3 rotation-xor operations at steps t_3 and t_6 can be interpreted as polynomial multiplications in the binary polynomial ring $\mathbb{Z}_2[x]/(x^{32} + 1)$:

$$t_3 = p * t_2 \mod x^{32} + 1, \text{ with } p = x^{25} + x^{24} + 1 \quad (7)$$

$$t_6 = q * t_4 \mod x^{32} + 1, \text{ with } q = x^{17} + x^9 + 1. \quad (8)$$

The inverse polynomials have Hamming weight 9:

$$p * (x^{28} + x^{21} + x^{20} + x^{14} + x^{12} + x^7 + x^6 + x^5 + x^4) \equiv 1 \pmod{x^{32} + 1} \quad (9)$$

$$q * (x^{27} + x^{19} + x^{18} + x^{17} + x^{11} + x^9 + x^3 + x^2 + 1) \equiv 1 \pmod{x^{32} + 1} \quad (10)$$

The choice of p and q guarantees that input (differentials) of weight less than 6 at t_2 and t_5 will always have output weight of at least 3 at t_3 and t_6 . Ignoring the nonlinear operation at step t_5 , the composite $p*q$ also has this property, but with guaranteed output weight of 4. The coefficients were chosen in a way to allow for a reasonably efficient implementation on AVR, which only has instructions for single-bit shifts of bytes.

There are some potentially problematic 4-bit to 4-bit rotational differentials such as $0x80808080 \lll$, but we could not cancel out the strong feedback propagation in our attack (with this particular p and q selection), which made exploitation difficult.

Round constants. The round constants of Equation 3 are just bytes from a maximum distance separable (MDS) code in decreasing-increasing order. The Hamming distance between each pair is at least 4. Efficient digital circuits can be constructed to generate this code – the last 8 bytes are just logical inverses of the first 8, for example. The modes described in this document only use the first 8 so implementations may choose not to include the last 8.

Doubling the number of rounds. Table 3 defines sixteen round constants as we reserve the option of doubling the number of rounds to $\rho = 12/16$ for extra margin of security.

The current ρ choices are based on quite optimistic estimates from avalanche and simple differential cryptanalysis. We encourage developers to choose the round-doubled versions for applications where throughput is not the main selection criteria (e.g. when hashing is only required only verifying signatures of firmware updates). A notation such as SNEIKEN- k - ρ and SNEIKHA- c - ρ may be used for these variants.

Note that schemes such as ChaCha are used with a wide array of different round selections [Ber08]. The Xoofff proposal uses a six-round Xoodoo [DHAK18], which is known to be vulnerable to algebraic distinguishers (Xoodoo has only been proposed as a secure permutation with 12 rounds).

Table 3: Performance comparison of some primitives. For Sponge permutations we give cycles/byte estimates for 128-bit and 256-bit capacity, corresponding to the security of an AEAD or square of security of a hash. Note that the RAM usage is not uniformly reported in the literature; clearly one needs RAM for the permutation in Sponge modes and for expanded keys in case of AES. We are reporting just the stack usage in this table.

Algorithm		ROM	Stack	– Cycles / Byte –		
				Round	128-bit	256-bit
Atmel AVR ATmega						
SNEIK	Fast [This work]	1974	14	16.8	135	270
SNEIK	Small [This work]	618	19	17.6	141	282
AES	Fast [Poe07]	3411	?	15.5	155	
AES	Small [Poe07]	1570	?	17.1	171	
Gimli	Fast [BKL ⁺ 17]	19218	45	8.88	320	639
Gimli	Small [BKL ⁺ 17]	778	44	17.2	620	1239
ARM Cortex M3/M4						
SNEIK	Fast [This work]	560	16	2.94	23.5	47.0
SNEIK	Small [This work]	232	28	3.31	26.4	52.0
Gimli	[BKL ⁺ 17]	3972	44	0.875	31.5	63.0
AES	Unprotected CTR [SS16]	2192/2960	72	≈ 3.5	34.7	49.5

Sponge modes. The BLNK2 modes are based on Author’s BLINKER framework for lightweight Sponge-based protocols [Saa14a], which has inspired derivative works such as Mike Hamburg’s STROBE [Ham17]. The mode implementation is derived from the one used for CBEAM [Saa14b] and WHIRLBOB [SB15] proposals.

We use an updated variant with a full-state keying mechanism and also a full-state keyed sponge method for associated data [GPT15, MRV15]. This full-state use case motivated us to move domain separation from capacity to be an “out-of-band” parameter of the cryptographic permutation itself. The capacity of SNEIKEN capacity matches the intended security level, as discussed in [JLM14].

Comparison to other schemes. Table 3 gives a performance comparison against some other candidates. Reliable and actually comparable data are difficult to find – for example AES is not directly comparable since there is no mode that satisfies the “ 2^{50} data under a single key” requirement of the NIST call for lightweight proposals. However, it is clear that SNEIK has quite advantageous performance and code size characteristics.

References

- [BBF⁺19] Hayo Baan, Sauvik Bhattacharya, Scott Fluhrer, Oscar Garcia-Morchon, Thijs Laarhoven, Ronald Rietman, Markku-Juhani O. Saarinen, Ludo Tolhuizen, and Zhenfei Zhang. Round5: Compact and fast post-quantum public-key encryption. In *PQCrypto 2019 – The Tenth International Conference on Post-Quantum Cryptography. Chongqing, China, May 8-10, 2019*, volume to appear of *Lecture Notes in Computer Science*. Springer, 2019. URL: <https://eprint.iacr.org/2019/090>.

- [Ben14] Josh Benaloh, editor. *Topics in Cryptology - CT-RSA 2014 - The Cryptographer's Track at the RSA Conference 2014, San Francisco, CA, USA, February 25-28, 2014. Proceedings*, volume 8366 of *Lecture Notes in Computer Science*. Springer, 2014. doi:[10.1007/978-3-319-04852-9](https://doi.org/10.1007/978-3-319-04852-9).
- [Ber08] Daniel J. Bernstein. Chacha, a variant of salsa20, 2008. URL: <https://cr.yp.to/chacha/chacha-20080128.pdf>.
- [BFM⁺18] Joppe W. Bos, Simon Friedberger, Marco Martinoli, Elisabeth Oswald, and Martijn Stam. Fly, you fool! faster frodo for the ARM cortex-m4. *IACR Cryptology ePrint Archive*, 2018:1116, 2018. URL: <https://eprint.iacr.org/2018/1116>.
- [BKL⁺17] Daniel J. Bernstein, Stefan Kölbl, Stefan Lucks, Pedro Maat Costa Massolino, Florian Mendel, Kashif Nawaz, Tobias Schneider, Peter Schwabe, François-Xavier Standaert, Yosuke Todo, and Benoît Viguier. Gimli : A cross-platform permutation. In Wieland Fischer and Naofumi Homma, editors, *Cryptographic Hardware and Embedded Systems - CHES 2017 - 19th International Conference, Taipei, Taiwan, September 25-28, 2017, Proceedings*, volume 10529 of *Lecture Notes in Computer Science*, pages 299–320. Springer, 2017. doi:[10.1007/978-3-319-66787-4_15](https://doi.org/10.1007/978-3-319-66787-4_15).
- [BSS⁺13] Ray Beaulieu, Douglas Shors, Jason Smith, Stefan Treatman-Clark, Bryan Weeks, and Louis Wingers. The SIMON and SPECK families of lightweight block ciphers. *IACR Cryptology ePrint Archive*, 2013:404, 2013. URL: <https://eprint.iacr.org/2013/404>.
- [DEMS16] Christoph Dobraunig, Maria Eichlseder, Florian Mendel, and Martin Schl  ffer. Ascon v1.2. *Submission to the CAESAR Competition*, 2016. URL: <https://competitions.cr.yp.to/round3/asconv12.pdf>.
- [DHAK18] Joan Daemen, Seth Hoffert, Gilles Van Assche, and Ronny Van Keer. The design of xoodoo and xooff. *IACR Trans. Symmetric Cryptol.*, 2018(4):1–38, 2018. URL: <https://doi.org/10.13154/tosc.v2018.i4.1-38>, doi:[10.13154/tosc.v2018.i4.1-38](https://doi.org/10.13154/tosc.v2018.i4.1-38).
- [GPT15] Peter Gazi, Krzysztof Pietrzak, and Stefano Tessaro. The exact PRF security of truncation: Tight bounds for keyed sponges and truncated CBC. In Rosario Gennaro and Matthew Robshaw, editors, *Advances in Cryptology - CRYPTO 2015 - 35th Annual Cryptology Conference, Santa Barbara, CA, USA, August 16-20, 2015, Proceedings, Part I*, volume 9215 of *Lecture Notes in Computer Science*, pages 368–387. Springer, 2015. doi:[10.1007/978-3-662-47989-6_18](https://doi.org/10.1007/978-3-662-47989-6_18).
- [Ham17] Mike Hamburg. The STROBE protocol framework. *IACR Cryptology ePrint Archive*, 2017:3, 2017. URL: <http://eprint.iacr.org/2017/003>.
- [JLM14] Philipp Jovanovic, Atul Luykx, and Bart Mennink. Beyond 2 c/2 security in sponge-based authenticated encryption modes. In Palash Sarkar and Tetsu Iwata, editors, *Advances in Cryptology - ASIACRYPT 2014 - 20th International Conference on the Theory and Application of Cryptology and Information Security, Kaoshiung, Taiwan, R.O.C., December 7-11, 2014. Proceedings, Part I*, volume 8873 of *Lecture Notes in Computer Science*, pages 85–104. Springer, 2014. doi:[10.1007/978-3-662-45611-8_5](https://doi.org/10.1007/978-3-662-45611-8_5).

- [KN10] Dmitry Khovratovich and Ivica Nikolic. Rotational cryptanalysis of ARX. In Seokhie Hong and Tetsu Iwata, editors, *Fast Software Encryption, 17th International Workshop, FSE 2010, Seoul, Korea, February 7-10, 2010, Revised Selected Papers*, volume 6147 of *Lecture Notes in Computer Science*, pages 333–346. Springer, 2010. URL: https://doi.org/10.1007/978-3-642-13858-4_19, doi:10.1007/978-3-642-13858-4_19.
- [MKJR16] Kathleen M. Moriarty, Burt Kaliski, Jakob Jonsson, and Andreas Rusch. PKCS #1: RSA cryptography specifications version 2.2. *RFC*, 8017:1–78, 2016. doi:10.17487/RFC8017.
- [MRV15] Bart Mennink, Reza Reyhanitabar, and Damian Vizár. Security of full-state keyed sponge and duplex: Applications to authenticated encryption. In Tetsu Iwata and Jung Hee Cheon, editors, *Advances in Cryptology - ASIACRYPT 2015 - 21st International Conference on the Theory and Application of Cryptology and Information Security, Auckland, New Zealand, November 29 - December 3, 2015, Proceedings, Part II*, volume 9453 of *Lecture Notes in Computer Science*, pages 465–489. Springer, 2015. doi:10.1007/978-3-662-48800-3_19.
- [Poe07] B. Poettering. Avraes: The aes block cipher on avr controllers, 2007. URL: <http://point-at-infinity.org/avraes/>.
- [Saa14a] Markku-Juhani O. Saarinen. Beyond modes: Building a secure record protocol from a cryptographic sponge permutation. In Benaloh [Ben14], pages 270–285. doi:10.1007/978-3-319-04852-9_14.
- [Saa14b] Markku-Juhani O. Saarinen. CBEAM: efficient authenticated encryption from feebly one-way ϕ functions. In Benaloh [Ben14], pages 251–269. doi:10.1007/978-3-319-04852-9_13.
- [SB15] Markku-Juhani O. Saarinen and Billy Bob Brumley. Whirlbob, the whirlpool based variant of STRIBOB. In Sonja Buchegger and Mads Dam, editors, *Secure IT Systems, 20th Nordic Conference, NordSec 2015, Stockholm, Sweden, October 19-21, 2015, Proceedings*, volume 9417 of *Lecture Notes in Computer Science*, pages 106–122. Springer, 2015. doi:10.1007/978-3-319-26502-5_8.
- [SS16] Peter Schwabe and Ko Stoffelen. All the AES you need on cortex-m3 and M4. In Roberto Avanzi and Howard M. Heys, editors, *Selected Areas in Cryptography - SAC 2016 - 23rd International Conference, St. John's, NL, Canada, August 10-12, 2016, Revised Selected Papers*, volume 10532 of *Lecture Notes in Computer Science*, pages 180–194. Springer, 2016. URL: https://doi.org/10.1007/978-3-319-69453-5_10, doi:10.1007/978-3-319-69453-5_10.
- [WN98] David J. Wheeler and Roger M. Needham. Correction to xtea. *Informal Manuscript or Report*, 1998. URL: <https://www.mjos.fi/doc/misc/xxtea.pdf>.